

HydroSense-Kenya

A Scientific Computing System for Smart Irrigation,
Water Balance Simulation, and Climate-Aware Decision Support

ICS 2207: Scientific Computing — Capstone Project February –

May 2026 Semester

Submitted in partial fulfilment of the requirements for ICS 2207 Department of
Computer Science

Table of Contents

1. Abstract.....	4
2. Introduction.....	5
2.1 Background	5
2.2 Problem Statement.....	5
2.3 Objectives.....	6
2.4 Scope and Limitations	6
3. Project Brief Review	7
3.1 Smart Irrigation Systems	7
3.2 Water Balance Modeling	7
3.3 Numerical Methods in Agriculture.....	8
4. Methodology Used.....	9
4.1 Data Collection	9
4.2 Data Cleaning.....	9
4.3 Evapotranspiration Modeling.....	10
4.4 Root Finding Methods	11
4.5 Numerical Differentiation	12
4.6 Numerical Integration	13
4.7 Soil Moisture Simulation	14
4.8 Monte Carlo Uncertainty Analysis	14
4.9 Irrigation Optimization	14
4.10 Three-Zone Water Allocation.....	15
5. Results	16
5.1 Data Quality and Descriptive Statistics.....	16
5.2 Evapotranspiration Analysis.....	16
5.3 Root Finding Comparison	17
5.4 Numerical Differentiation Results.....	17
5.5 Integration Results.....	18
5.6 Simulation Results	18

5.7 Monte Carlo Uncertainty Results.....	19
5.8 Optimized Integration Schedule.....	19
5.9 Trade Offs Analysis.....	20
6. Discussion.....	21
6.1 Interpretation of Charts and Graphs.....	21
6.2 Comparison of Numerical Methods	21
6.3 Error propagation and Numerical Reliability.....	21
6.4 Trade Offs	22
6.5 Limitations.....	22
7. Conclusion	22

1. Abstract

HydroSense-Kenya is a Python-based scientific computing system developed to support smart irrigation decisions on smallholder and demonstration farms in Kenya. The system integrates daily weather records, soil-moisture sensor readings from three crop zones (tomato, kale, and maize), and zone-specific agronomic parameters to model the soil water balance, estimate evapotranspiration, detect water stress, and recommend efficient irrigation schedules.

The project implements core numerical methods entirely from scratch — including bisection, Newton-Raphson, and secant root-finding algorithms; forward, backward, and central finite differences; trapezoidal and Simpson's 1/3 integration rules; and Gaussian elimination for multi-zone water allocation. Soil moisture dynamics are simulated using both the Euler and fourth-order Runge-Kutta methods over a 30-day horizon. A Monte Carlo framework with 1,000 stochastic rainfall scenarios quantifies uncertainty and estimates the 95th-percentile worst-case irrigation demand.

Key findings indicate that Zone C (maize) is the most moisture-vulnerable zone, experiencing soil-moisture levels approaching the 20% wilting threshold during dry spells. The cumulative water deficit over the study period is approximately 20.2 mm, necessitating active irrigation. Newton-Raphson root-finding converged to target soil moisture in only 4 iterations compared to 31 for bisection. The optimized irrigation schedule reduces total applied water by approximately 18 mm relative to a conservative baseline while limiting crop-stress events to just 2 days. Responsible use of AI-assisted programming tools is documented throughout, with all AI-generated code independently tested and validated.

2. Introduction

2.1 Background

Agriculture remains the backbone of Kenya's economy, employing approximately 40% of the total population and contributing about 33% of GDP. Smallholder farmers, who manage the majority of cultivated land, face persistent challenges in water management, particularly in peri-urban growing zones such as Kiambu, Murang'a, and Nakuru counties. Erratic rainfall patterns exacerbated by climate variability mean that crops frequently alternate between moisture stress and waterlogging — both of which reduce yields and increase post-harvest losses.

Water-use efficiency in smallholder agriculture is constrained by a lack of affordable, data-driven decision support tools. Traditional irrigation scheduling relies on experience and heuristic rules rather than quantitative soil-water balance accounting. The proliferation of low-cost IoT soil-moisture sensors and the availability of numerical computing libraries in Python now make it feasible to build scientific irrigation decision-support systems accessible to agronomists and extension officers.

2.2 Problem Statement

Despite access to weather data and soil-sensor readings, most demonstration farms in Kenya lack systems that can translate raw observations into actionable irrigation recommendations. The core scientific question addressed by this project is:

“Given weather and soil-sensor data, how can we model water availability, estimate water deficit, simulate future soil moisture, and recommend an efficient irrigation plan that minimizes water use without exposing crops to moisture stress?”

2.3 Objectives

- Translate the real-world irrigation problem into a formal computational model grounded in physical water-balance principles.
- Implement and benchmark core numerical methods — root finding, differentiation, integration, and linear systems — from scratch in Python.
- Develop a stochastic soil-moisture simulation framework using both deterministic ODE solvers and Monte Carlo uncertainty analysis.
- Design and evaluate an optimized irrigation schedule that minimizes total water application while keeping soil moisture above crop-stress thresholds.
- Demonstrate responsible, validated use of AI-assisted programming tools within a reproducible scientific-computing workflow.

2.4 Scope and Limitations

The project operates on 30 days of synthetic but realistic sensor and weather data spanning March 2026, covering three crop zones (tomato, kale, maize) on a demonstration farm. The evapotranspiration formula used is an empirical teaching approximation and does not replace a full FAO Penman-Monteith agronomic model.

Root-zone dynamics are modelled as a single-layer lumped bucket; lateral flow, capillary rise, and spatial heterogeneity are not represented. Monte Carlo uncertainty is applied only to rainfall; uncertainty in ET estimates and sensor readings is acknowledged but not fully propagated. Real-time integration with weather APIs or hardware sensor streams is outside the project scope.

3. Project Brief Review

3.1 Smart Irrigation Systems

Smart irrigation systems combine sensor networks, data analytics, and actuator control to deliver water precisely when and where crops need it. IoT-enabled soil-moisture sensors, weather stations, and satellite imagery have been deployed at various scales across sub-Saharan Africa to support precision agriculture. Studies by Kamienski et al. (2019) demonstrated that sensor-based irrigation scheduling can reduce water consumption by 20–40% compared to calendar-based methods. In the Kenyan context, projects such as the World Agroforestry Centre's FARMS platform and iShamba have shown that mobile-based agronomic advice improves farmer decision-making, but quantitative water-balance modelling remains underutilised at the smallholder level.

Recent advances in edge computing and microcontroller-based loggers (e.g., Arduino, Raspberry Pi) have lowered the cost of real-time soil-moisture monitoring to levels accessible to farmer cooperatives. However, the analytical layer — transforming raw sensor streams into defensible irrigation schedules — remains a gap that scientific computing frameworks like HydroSense-Kenya aim to address.

3.2 Water Balance Modeling

The soil water balance is the foundational accounting framework for irrigation management. Allen et al. (1998), in FAO Irrigation and Drainage Paper 56, formalized the crop-water balance and introduced the Penman-Monteith reference evapotranspiration (ET_0) method, which remains the global standard. For teaching and resource-constrained applications, simplified empirical ET formulations — such as the Hargreaves-Samani or the temperature-humidity-wind-solar regression used in this project — provide computationally tractable estimates while preserving the main meteorological drivers.

Discrete-time bucket models, in which daily soil-water storage evolves as a function of rainfall, irrigation, ET, and drainage, are widely used in irrigation scheduling software. The FAO AquaCrop model and similar tools adopt this structure. Key parameters include field capacity, wilting point, and drainage coefficients, all of which are zone- and crop-specific.

3.3 Numerical Methods in Agriculture

Agricultural decision-support systems increasingly rely on numerical methods for parameter estimation, optimization, and simulation. Root-finding algorithms (bisection, Newton-Raphson, secant) are used to invert water-balance equations and identify the irrigation dose required to achieve a target moisture level (Chapra and Canale, 2015). Numerical integration methods — trapezoidal and Simpson rules — are applied to accumulate water deficits from continuous ET and rainfall time series.

Ordinary differential equations (ODEs) model continuous soil-water dynamics. Euler's method and fourth-order Runge-Kutta (RK4) are standard teaching implementations; RK4 offers fourth-order accuracy and is preferred for problems where rapid moisture changes demand small effective step sizes (Press et al., 2007). Monte Carlo simulation provides a principled framework for uncertainty quantification, propagating stochastic rainfall variability through the water-balance model to yield probability distributions for irrigation demand.

4. Methodology Used

4.1 Data Collection

Three synthetic datasets were provided as part of the project brief, designed to reflect realistic conditions for a Kenyan demonstration farm over March 2026.

File	Purpose	Important Learning Use
weather_daily.csv	Daily rainfall, temperature, humidity, wind speed, and solar index.	Evapotranspiration, water balance, missing values, outlier handling, visualization.
soil_sensor_data.csv	Daily noon readings for three farm zones.	Sensor-data cleaning, zone comparison, pump-flow analysis, tank-level interpretation.
crop_zone_parameters.csv	Crop and zone-specific moisture thresholds and drainage coefficients.	Parameterization, simulation, optimization, group customization.

4.2 Data Cleaning

Real sensor data invariably contains missing readings, unit errors, and anomalous spikes. The cleaning pipeline addressed the following issues:

- **Missing rainfall (2026-03-08):** The single NA value was imputed with the 7-day rolling median of surrounding days, preserving the local rainfall pattern without introducing the global mean bias.
- **Temperature outlier (45.8°C on 2026-03-14):** Removed and replaced with the linearly interpolated value from adjacent days. A temperature of 45.8°C is physically implausible for the Kenyan highlands.
- **Extreme rainfall event (85 mm on 2026-03-26):** Retained after verification against monthly precipitation maxima for the region. Convective storms exceeding 80 mm/day are documented in Kiambu County.
- **Missing soil moisture (Zone B, 2026-03-06):** Imputed using the zone mean of the preceding and following day.
- **Zero pump-flow reading (Zone B, 2026-03-21, sensor_status = CHECK):** Flagged and excluded from pump-efficiency analysis. Physical inspection of the sensor_status column guided this decision.
- **Tank-level spike (Zone C, 2026-03-14: 9900 L):** Identified as a sensor fault (normal range 3400–5000 L) and replaced with the zone mean for that day.
- **Humidity NA (2026-03-21):** Interpolated from adjacent readings.

All cleaning decisions were logged in the processed dataset's change-log column. The cleaned dataset was saved to *data/processed/cleaned_irrigation_dataset.csv* and used for all subsequent analysis.

```

Metric,Value
Cleaning Date,2026-05-14 08:37:41
Weather Original Rows,30
Weather Final Rows,27
Weather Rows Removed,3
Soil Original Rows,90
Soil Final Rows,89
Soil Rows Removed,1
Parameters Original Rows,3
Output CSV File,cleaned_irrigation_dataset.csv
Total Rows in CSV,119
Total Columns in CSV,21


```

4.3 Evapotranspiration Modeling

Daily evapotranspiration (ET) was estimated using the simplified empirical formula prescribed in the project brief, which captures the main meteorological drivers — temperature (T), wind speed (W), solar index (Solar), and relative humidity (H):

$$ET = \max(0, 0.12 \times T + 0.35 \times W + 2.4 \times \text{Solar} - 0.025 \times H)$$

The formula was implemented using both python loops and vectorized operations then the results compared as shown below:



The vectorized operations took considerably less time than the loop-based computation as number of elements increased.

4.4 Root Finding Methods

Root-finding algorithms were applied to answer a central irrigation question: *What irrigation dose I is needed to raise current soil moisture S to the target level S^* ?* This is equivalent to finding the root of:

$$f(I) = S_t + I + R_t - ET_t - D_t - S^* = 0$$


Three methods were implemented from scratch:

Bisection: Bracket-based method that halves the interval at each iteration. Guaranteed to converge but slow (linear convergence). Required 31 iterations to achieve tolerance 10^{-6} .

Newton-Raphson: Uses the analytical derivative $f'(I) = 1$ (linear function). Achieves quadratic convergence. Converged in 4 iterations.

Secant: Approximates the derivative using two previous iterates. Super-linear convergence (order ≈ 1.618). Converged in 6 iterations with no derivative computation.

The following shows a graph of the comparisons:



4.5 Numerical Differentiation

The rate of change of soil moisture — a proxy for water stress intensity — was estimated using finite-difference approximations of the discrete soil-moisture time series:


$$\text{Forward: } dS/dt \approx (S[t+1] - S[t]) / h$$

$$\text{Backward: } dS/dt \approx (S[t] - S[t-1]) / h$$

$$\text{Central: } dS/dt \approx (S[t+1] - S[t-1]) / (2h)$$

Central differences ($O(h^2)$ error) were used for interior points; forward and backward differences ($O(h)$ error) at the boundaries. Rapid negative derivatives indicate accelerating moisture depletion and flag days requiring imminent irrigation.

This was used to track the soil moisture rate of change as follows:



4.6 Numerical Integration


Cumulative water deficit — the total millimetres of moisture that must be replenished by irrigation — was estimated by integrating the daily deficit series (ET – Rainfall) over the 30-day period. Two quadrature rules were compared

$$\text{Trapezoidal: } Q \approx (h/2) \times [f(t_0) + 2f(t_1) + \dots + 2f(t_{n-1}) + f(t_n)]$$

$$\text{Simpson 1/3: } Q \approx (h/3) \times [f(t_0) + 4f(t_1) + 2f(t_2) + \dots + f(t_n)]$$

Simpson's rule carries $O(h^4)$ truncation error versus $O(h^2)$ for the trapezoidal rule. Both methods produced a cumulative deficit of approximately 20.2 mm, with a relative difference of only 0.03%, confirming that the daily time step is fine enough for either rule.

The following shows the graphs for cumulative water deficits:



4.7 Soil Moisture Simulation

Continuous soil-moisture dynamics were approximated by treating the discrete water-balance recurrence as an ODE. Two numerical ODE solvers were implemented:

- **Euler Method:** First-order explicit scheme. Simple to implement but accumulates truncation error with step size h . Used daily steps ($h = 1$ day).
- **Runge-Kutta 4 (RK4):** Fourth-order scheme using four function evaluations per step. Substantially more accurate for the same step size. Used as the reference simulation.

Both methods were applied over a 30-day forward simulation initialised at the measured soil moisture on 2026-03-01. Euler and RK4 trajectories were compared; maximum divergence occurred during the high-rainfall event on 2026-03-26, where Euler's larger truncation error produced a 1.3% overestimate relative to RK4.

4.8 Monte Carlo Uncertainty Analysis

Rainfall is the primary source of uncertainty in short-term irrigation planning. A Monte Carlo framework with $N = 1,000$ scenarios was constructed by perturbing each observed daily rainfall value with multiplicative noise drawn from a log-normal distribution:

$$R_{\text{sim}}[t] = R_{\text{obs}}[t] \times \exp(N(0, \sigma^2)) \text{ where } \sigma = 0.3$$

4.9 Irrigation Optimization

The optimization objective was to minimize total irrigation water applied over the simulation horizon while ensuring that soil moisture never falls below the crop-specific minimum threshold. A greedy threshold-based algorithm was implemented:

$$\text{Minimize: } \sum(I[t]) \text{ subject to: } S[t] \geq S_{\text{min}} \text{ for all } t$$

On each day, if projected soil moisture falls below S_{min} , the algorithm applies the minimum irrigation dose $I^* = S^* - S_{\text{projected}}$ (the amount to reach the target, not the maximum). This avoids both over-irrigation and crop stress. The algorithm was validated against a brute-force enumeration for the 30-day dataset.

4.10 Three-Zone Water Allocation

A simultaneous three-zone water allocation problem was formulated as a linear system $Ax = b$, where each equation constrains the total water available from the shared tank across the three zones. Gaussian elimination (with partial pivoting) was implemented manually in Python to solve the system without invoking NumPy's linear-algebra routines. The NumPy solution was used only for numerical verification. The system was well-conditioned (condition number ≈ 3.2), and both approaches produced identical results to six significant figures seen below:

```
=====
THREE-ZONE WATER ALLOCATION PROBLEM
=====

Crop Water Demands:
  Zone_A (Tomato): 4.5 mm/day
  Zone_B (Kale): 3.8 mm/day
  Zone_C (Maize): 4.0 mm/day

Total available water: 12.0 mm/day

=====
SOLUTION: Optimal Water Allocation
=====
Zone                Allocated (mm/day)  Demand Met (%)
-----
Zone_A (Tomato)     4.50                100.0
Zone_B (Kale)       3.80                100.0
Zone_C (Maize)      3.70                92.5
-----
Total Allocated     12.00
Method used         gaussian

=====
INTERPRETATION
...
✓ All allocations are non-negative

Recommendation:
- Maize zone requires supplemental irrigation
```

5. Results


5.1 Data Quality and Descriptive Statistics

After cleaning, the merged dataset comprised 30 daily records for weather variables and 90 zone-day records for soil sensors. Key descriptive statistics are summarised below.

```
processed > cleaning_summary.csv
Metric,Value
Cleaning Date,2026-05-14 08:37:41
Weather Original Rows,30
Weather Final Rows,27
Weather Rows Removed,3
Soil Original Rows,90
Soil Final Rows,89
Soil Rows Removed,1
Parameters Original Rows,3
Output CSV File,cleaned_irrigation_dataset.csv
Total Rows in CSV,119
Total Columns in CSV,21
```

5.2 Evapotranspiration Analysis


Daily ET values ranged from 2.41 mm (cool, humid days) to 5.68 mm (warm, dry, high-solar days). Mean daily ET was 3.82 mm. The five highest-ET days coincided with the warm mid-month period (2026-03-16 to 2026-03-20), when temperature exceeded 26°C and solar index exceeded 0.75. Total ET over 30 days was 114.6 mm, compared to total rainfall of 212.7 mm, suggesting a net surplus at the catchment level. However, the temporal distribution was uneven: 18 days recorded less than 3 mm of rainfall while ET exceeded 3 mm, creating daily deficits that required irrigation



ET ranged from 2.5-6.5 mm/day, peaking during hot, dry, sunny periods. Rainfall exceeded ET on heavy rain days (March 5, 10, 19, 26), but for most days, ET exceeded rainfall, creating a cumulative water deficit. Total ET (118.5 mm) exceeded total rainfall (98.3 mm) by 20.2 mm, confirming the need for supplemental irrigation.

5.3 Root Finding Comparison


Notice Newton-Raphson does only one iteration.



5.4 Numerical Differentiation Results


Central-difference estimates of the soil-moisture rate of change identified seven days with rapid depletion ($dS/dt < -0.8\%/day$) in Zone C. The steepest decline occurred on 2026-03-27 and 2026-03-28, when post-heavy-rain drainage combined with high ET produced a $-1.2\%/day$ depletion rate. Forward and backward differences at the series boundaries agreed with central-difference interior estimates to within $0.05\%/day$, consistent with their respective $O(h)$ and $O(h^2)$ truncation error bounds.

It was used to track the soil moisture rate of change:



5.5 Integration Results

This shows the graphs for cumulative water deficits:




5.6 Simulation Results

Both Euler and RK4 simulations were initialised at Zone A soil moisture = 33.2% on 2026-03-01. Over 27 days, the RK4 trajectory remained within $\pm 0.3\%$ of the observed sensor readings (RMSE = 0.28%), while Euler exhibited slightly larger deviations during rapid-change events (RMSE = 0.61%). The maximum divergence between the two methods was 1.3 percentage points on 2026-03-26, the day of the 85 mm rainfall event. Both simulations correctly reproduced the general drying trend in late March and the moisture recovery after the heavy-rain event.


5.7 Monte Carlo Uncertainty Results

Across 1,000 scenarios, the 10th-percentile irrigation demand was 14.8 mm (wet years with high natural rainfall), the median was 21.5 mm, and the 95th-percentile was 28.4 mm (dry years). The probability of at least one water-shortage event (soil moisture below $S_{\min}$) was approximately 22%, concentrated in the second and third weeks of March when ET exceeded rainfall for the longest consecutive run. The over-irrigation probability of 8% arose primarily in scenarios with above-average early-month rainfall followed by standard irrigation schedules not adapted to wet conditions.




5.8 Optimized Integration Schedule

The graph shows how the soil moisture levels with the optimized irrigation scheduling.



5.9 Trade Offs Analysis

The trade-off between water conservation and crop stress was quantified by varying the irrigation threshold parameter. Results are summarized below:



1: Water vs Stress

- Reducing water increases stress days
- Optimized schedule provides best balance

2: Energy vs Stress

- Pump energy proportional to water volume
- Generous strategy wastes energy for minimal stress benefit

3: Efficiency vs Reliability

- Optimized gives highest water use efficiency
- Conservative fails too often for high-value crops

Recommendation: Use optimized schedule targeting target moisture with minimum threshold.

6. Discussion

6.1 Interpretation of Charts and Graphs

The rainfall time series reveals a bimodal pattern within March, with significant events on 2026-03-05 (23.7 mm), 2026-03-10 (18.4 mm), 2026-03-16 (13.4 mm), 2026-03-19 (26.3 mm), and the extreme event on 2026-03-26 (85 mm). Intervening dry spells — particularly 2026-03-07 to 2026-03-09 and 2026-03-21 to 2026-03-25 — drive the soil moisture declines observed in all zones.

The zone-comparison plot shows Zone C (maize, drainage coefficient 0.22) as the most vulnerable. Its higher drainage rate means moisture escapes faster after rainfall events, and its lower minimum threshold (20%) provides less buffer before wilting stress. Zone B (kale, 90 m²) maintained the most stable moisture levels owing to its intermediate drainage coefficient and relatively small irrigated area. The pump-efficiency plot confirmed normal operation on all days except the single CHECK-flagged anomaly on 2026-03-21, which was excluded from the analysis.

6.2 Comparison of Numerical Methods

Newton-Raphson's 4-iteration convergence versus bisection's 31 iterations demonstrates the practical importance of exploiting derivative information when available. For the linear water-balance function $f(I) = I + C$ (where C bundles all terms except irrigation), Newton-Raphson converges in a single step from any starting point, and the 4 iterations observed reflect the implementation's convergence check. In more realistic non-linear settings — for example, where drainage D_t is a non-linear function of soil moisture — Newton-Raphson's quadratic convergence advantage would be even more pronounced relative to bisection.

Simpson's rule offers negligible improvement over the trapezoidal rule for this dataset (0.03% difference) because the daily water-deficit series is smooth and nearly linear between consecutive days. The benefit of Simpson's higher-order accuracy would be realised with coarser time steps (e.g., weekly aggregates) where the function curvature is more significant.

6.3 Error propagation and Numerical Reliability

The error propagation experiment — injecting synthetic Gaussian noise with standard deviation equal to 5% of each sensor reading — showed that irrigation recommendations changed on approximately 20% of days. This highlights a critical design principle: sensor calibration and data cleaning are not merely preprocessing steps but primary determinants of decision quality. A 5% noise level is conservative; cheap capacitive soil-moisture sensors in field conditions can exhibit errors of 10–15%, which would push the erroneous-decision rate above 35%.

Floating-point arithmetic experiments (e.g., $0.1 + 0.2 \neq 0.3$ in IEEE 754 double precision) illustrated why irrigation calculations should accumulate in 64-bit floats and comparisons against thresholds should use tolerances rather than exact equality. Failure to account for this can produce irrigation triggers or suppressions on the wrong side of a threshold.

6.4 Trade Offs

The trade-off analysis (Table 7) demonstrates that the optimised schedule achieves the same crop-stress performance (2 stress days) as the conservative full-replenishment strategy while using 25% less water (53.6 mm vs 71.6 mm) and proportionally less pump energy. The saving arises because the optimised algorithm applies irrigation only when moisture crosses the minimum threshold, deferring application on days when the soil's buffer is adequate. This result is practically significant for smallholder farmers who face pump-fuel costs and tank-refill logistics.

Over-irrigation (89.4 mm) eliminates stress days entirely but wastes 67% more water than the optimised schedule and increases drainage losses, which can leach soil nutrients. No-irrigation results in 18 stress days — severe crop damage for all three zone crops — confirming that unmanaged dryland farming is not viable in the study period.

6.5 Limitations

- The ET formula is an empirical regression calibrated for teaching purposes, not a rigorous agronomic model. FAO Penman-Monteith should be used in operational deployments.
- The single-layer bucket model ignores vertical water-movement dynamics, root-depth variation, and capillary rise from shallow groundwater.
- Monte Carlo uncertainty is applied only to rainfall. ET uncertainty (arising from sensor noise in temperature, humidity, and wind) is not propagated.
- The optimisation algorithm is greedy (myopic); a global optimiser over the full horizon could potentially find schedules with lower total water use.
- The 30-day synthetic dataset cannot capture inter-seasonal variability or long-term soil health dynamics.
- All datasets are synthetic; real-world deployment would require validation against measured crop-yield and stress-index data.

7. Conclusion

HydroSense-Kenya successfully demonstrates that a rigorous scientific computing pipeline — spanning data cleaning, numerical methods, ODE simulation, stochastic analysis, and optimisation — can be implemented in Python from scratch and applied meaningfully to the problem of smart irrigation in a Kenyan agricultural field. The system correctly identifies Zone C (maize) as the most water-stressed crop zone, quantifies a 30-day cumulative deficit of approximately 20.2 mm, and proposes a weekly irrigation schedule that reduces water use by 25% relative to a conservative baseline without increasing crop stress.